technology and the team knowledge quickly become scale problems. Eventually, more time is spent on maintenance, and new projects become more and more delayed.

The other benefit of data mesh is the clarity of data ownership right from the point of its creation. This approach flattens the data management structure, leaving just a thin layer of a federated governance team. And even that team's activities are limited to agreeing on standards within autonomous domains.

The speedup of development also comes from empowering the implementation teams. Since producing and maintaining the data products rests on their shoulders, the speed of change is not limited by a single central integrations team's backlog of tasks. Therefore, both the evolution of and fixes to the data product happen more quickly—especially any bug fixing and downtime. Furthermore, the team that owns the data product is better equipped to react faster because no context switching is needed, as in the case with one central team.

The other factor worth mentioning is data environment stability. With data products offering access to contracted versions of their datasets, pipelines built on them are much more robust and require much less maintenance.

1.3 Use case: A snow-shoveling business

Candace operates a snow-shoveling business. She is a proud businesswoman who started the business shoveling snow herself every winter. After a couple of years, she scaled up her operations. She focused on logistics, billing, and setting the prices across the business, and she hired three employees: Adam, who shovels the houses at Pine Road; Eve, who does the houses on Oak Street; and Bob, who is responsible for ordering new shovels, because they seem to break all the time.

Adam and Eve are responsible for both shoveling snow and bringing in new clients on their respective streets. After all, they spend quite some time around people there while shoveling snow each winter! But Candace isn't happy. Figure 1.4 depicts the initial state of Candace's operations.