- At least one of your domains has mostly structured data sources.
- You are extracting value from that domain's data by using BI, reporting, and data analysis.
- Your data engineers have expertise in data warehouse technologies.

2.2.2 Data lake

We mentioned data lakes in chapter 1. Here we will go into greater detail. A *data lake* is quite similar to a data warehouse; for example, both rely on a central data repository. The main difference is in the way a data lake stores data: in raw format. Data engineers transform the data into more structured and cleansed forms. They then move it into different zones along with its lineage and versioning.

A data lake has numerous *zones* (e.g., *process zones* used to store data already processed but not yet ready for production or accessed by end users; or *access zones*, with data already curated and ready to be accessed by consumers; see figure 2.4). A data lake is often described as a schema-on-read, in contrast to the data warehouse, which is described as a schema-on-write. In a data lake, incoming data is not forced to conform to the specific schema and data model. Instead, we force the data to conform before end users access it in different zones.

