community will find more interesting and beneficial ones. With every heuristic, we are trying to show you how it could be used.

Align data products with domains/business capabilities

This is the most important heuristic. It should be your default solution. As we showed you in sections 4.1 and 4.2, domain or business capability boundaries are very natural boundaries for one or more data products.

Align data products with possible usage or use cases

As we described in section 4.3, the other most widely applicable boundary is the group of related use cases. The best way to group use cases is to collate them by their purposes and business value. Data serving this group of use cases is a good candidate for data domain and data product boundaries.

Align data products with personas of consumers

Another heuristic related to the previous one is to group use cases based on the persona/actor/consumer performing the analysis or for whom those analyses are done.

The Movie Trends data product is an excellent example of how this heuristic can be applied. We established that we have two consumers: a marketing specialist and a content strategy lead. If we use this heuristic, we could create a Content-Producing Movie Trends data product and a Marketing Movie Trends data product. This setup has one advantage over a single Movie Trends data product: there is no risk of model ambiguity, where one entity means different things in marketing and movie production contexts. When we have two different consumers, the model could become too complex and hard to maintain. Although choosing heuristics usually requires a tradeoff, we choose between consistency and cost-efficiency (a single data product should be cheaper to maintain).

Align data products with the source system

In most cases, source systems are owned by teams that know the most about the data they produce. This is why it makes sense to create one or more data