We have analysts and other development teams as our *data consumers*. Again, we have two very different kinds of consumers with different requirements.

Finally, we have a lot of *data products*. Considering data products as a party is a big advantage. First of all, it displays the true nature of the interaction. Data consumers' prime interests are the pieces of data, not the people producing them. Ideally, that's all they would need, but of course, that is never enough because data consumers need context to work with data.

So the second interest of data consumers is to communicate with the data producers. In the data mesh, we aim to have that communication automatic because any personal communication does not scale into the mesh. Automatic communication happens primarily through the data catalog or the discovery endpoints of the data products or any other automatic mechanism we build into the self-serve data platform.

Communication through the data catalog is simple: data producers provide metadata, lots of it, and it's well maintained. Data consumers read and consume the metadata. That's it. That's all that should be needed. If that is not enough, the data products are likely not described well enough. This highlights that another goal of the self-serve data platform is to facilitate communication between data producers and data consumers.