The data platform team could have chosen to provide one hosted Airflow cluster maintained by the data platform team (or a third party like Amazon Web Services), with a reasonable separation between the compute environments per team.

Either option can be valid, depending on the circumstances. The question is simply, is the combination of platform team + flow-aligned teams able to increase their throughput with this option? And that includes answering whether maintaining a central service is feasible and whether flow-aligned teams (in this case, development teams) are able to manage things on their own.

We will now take a look at the final architecture of our self-serve data platform. Again, what the data platform team does is modularize and extend another part of the capabilities of the platform. Our new platform kernel contains three capabilities, pictured in figure 7.12:

- A data product plane for creating, finding, and discovering.
- A data storage plane for the internal creation of the default CSV storage, or the option to create a Terraform template for using a suitable PostgreSQL instance that will autoregister the data products' storage location.
- A data transformation plane can be used optionally. In this case, it creates and links a Terraform template for Airflow.