Figure 7.14 The final architecture adds yet another interface for the Airflow infrastructure.

The final architecture, shown in figure 7.14, now allows for each data-producing team to have decently equipped data products.

They now include the following:

- An optional CSV output port
- An optional SQL output port in the form of the SQL endpoint of a PostgreSQL instance
- An Airflow instance to easily manipulate data to prepare it for consumption
- The centralized metadata endpoints inside the DataHub

The data consumer stack consists of the DataHub and related tooling to mainly focus on finding and understanding the right data.

The workflow for data producers changes a lot with this newly provided interface. Both the Yellow and Gray teams decide to use the provided templates to start an Airflow instance inside their team context. They are using Airflow to quickly churn out new data products using their operational systems' REST APIs and event stores as input ports. Here are the details:

- Team Gray is now using just the SQL ports.
- Team Yellow opts to use the SQL endpoints in addition to the CSV port as output ports because the team is now free to create so many more valuable data products.
- Team Green keeps on using the CSV dumps, as it doesn't have the need for anything else right now.

The data consumers, both the ML engineers as well as analysts, use the central DataHub to locate datasets and then use their personal weapons to consume and combine them.

FEDERATED COMPUTATIONAL GOVERNANCE IN ITERATION 3