on the data consumer side.

8.2.4 Workflows

The data producer workflow in this technical architecture is similar to the one in the GCP section. Data producers pull the latest template from the central repository, configure it to meet their demands, opt in or out of possible Kinesis data stream creation, and apply the template to their AWS accounts.

The new addition is the PII service they can opt to use. If they want to use this service, they can integrate calls into the Glue code to the central service, which makes sure to hash personalized data equally.

The regular data consumer by default can use an open-for-all IAM role to access the metadata catalog. Then they use their own tools or the central Athena interface to query the related data. If they don’t have access to a certain data product, they can request it from the owning team. That team would then modify the IAM to provide access.

However, the data consumer can also opt to use the Terraform template to get a larger toolset. If they choose to, they need to configure it and then apply it to their account. Once done, they will have access to a ready and set SageMaker to explore data using Python inside a notebook. Once they are done with their analysis, they can also opt to create their own data products out of their analysis and publish them to the Lake Formation, following the same process the other teams are using.

8.2.5 Relation to data mesh ideas

Let’s discuss the data mesh principles indirectly used in this architecture.

Platform thinking, as a principle, is about making the lives of all platform participants easier, while not restricting their freedom. This architecture is aimed not just at data producers but also at data consumers, making it easy for a non-AWS savvy team to get going. It also goes slightly further than the Google Cloud version by allowing data producers to easily ingest arbitrary data sources.