In the AWS data mesh, data producers are pretty much forced into owning their data, because they also own the infrastructure. That will not be the case here.

On the flip side, this data mesh likely doesn't need more than one data engineer to maintain the technical infrastructure inside a small- to medium-sized company. Inside domains, you should be able to cover the needs with one data scientist/engineer per domain.

8.3.1 Self-serve data platform architecture

We're going to make a small modification to our teams. Teams Black and White are both development teams with one data scientist. Team Gray is just a one-person team maintaining the platform on its own. Our business analysts are SQL savvy and can do some work with Python inside notebooks.

Figure 8.9 depicts our architecture this time. Notice how the contexts are a bit more blurry. You could, of course, separate out the contexts of teams Black and White, but we chose not to in this case.