particular, it is an interesting setup for companies with data-intensive applications.

8.4 Data mesh on Kafka

Whereas the data mesh on Databricks described in the preceding section was effectively using a central technology piece for transforming data, but not for the initial collection, the architecture in this section starts at the initial collection of data.

With this, we're providing an even larger degree of centralization, which enables very fast building of data-intensive applications. Remember, we're talking about physical and technological centralization, not a logical one.

This setup will provide standardized input and output ports, without forcing Kafka onto a team. It will enable SQL-based transformation tooling and will bring the possibility to work nicely with legacy data sources that aren't yet migrated into the data mesh itself.

8.4.1 Self-serve data platform architecture

Our teams are equipped similarly to the setups in the first two sections. We have two software engineering teams, Black and White, and a platform team responsible for operating a few additional components that round up the Kafka-based data mesh.