- You got to know four architectures of self-serve platforms: a cloud-native architecture, a cloud-native architecture with larger autonomy for the data producers, a much easier-to-deploy architecture using a certain degree of centralization, and finally an architecture focused on building data-intensive real-time applications.
- We've covered a variety of toolkits—most importantly GCP, AWS, Databricks, and Apache Kafka—but also Terraform and DataHub.
- All of the architectures, as well as the toolkits, can be mixed and matched to solve your specific problems.
- Building a self-serve data platform is serious work simply because no appropriate out-of-the-box tooling is available.
- Because your underlying data landscape will likely be unique, you will end up customizing a solution for your company anyhow.
- The architectures and toolkits in this chapter will help you take a good stab at the first few iterations of building your own self-serve data platform, whether you start with a platform that needs almost no maintenance to maintain (like the Databricks-based example) or you want to build multiple real-time data-based applications (and thus opt for a Kafka-based setup).
- Whatever architecture you choose, you will end up facing a tradeoff of the data-producing teams' autonomy and flexibility against the cost and maintainability of the platform itself.