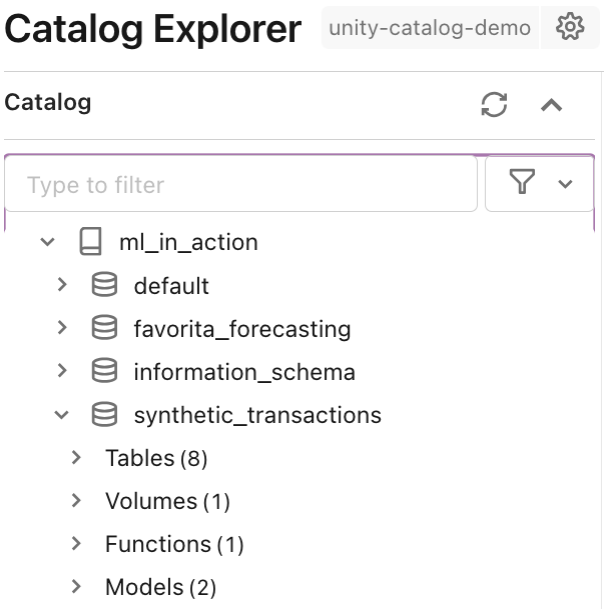


Figure 2.4 – UC governs all assets under catalogs, including databases, tables, volumes, functions, and models

Deciding on your metastore does not have to be a permanent choice. However, migrating later could become a headache. UC is continually improving and integrating with new Databricks features. The list of reasons to choose UC over HMS continues to grow, and our recommendation is to begin with and stick with UC.

To determine whether Unity Catalog is the right choice for you and your company, you can check out the shrinking list of limitations for choosing UC in *Further reading*. As UC continues to expand in capability, it is the path of the future. Specifically, for machine learning, there is an integration with the Feature Engineering client and the new Model Registry. Using the UC Model Registry for model sharing and governing is simpler. We will cover more about the Model Registry in Unity Catalog and Databricks Feature Engineering Client in *Chapters 5, 6*, and *7*, but if you're curious and eager to learn more now, you can check out Manage model lifecycle in Unity Catalog in the Further reading section. Given the ever-growing number of reasons to use UC, all project code in this book will use UC.

Defining where the data lives, and cloud object storage

All data products start with data, and so how we make the data accessible for data teams is another important early design choice. Databricks is a cloud-based platform that connects to cloud object storage – **Azure Data Lake Storage** (**ADLS**), Amazon **Simple Storage Service** (**S3**), or **Google Cloud Storage** (**GCS**). There is a separation of compute and storage. Databricks orchestrates compute; the data is in cloud object storage.