This is the ideal time to install the libraries you will need. Install via `PyPI opendatasets`, `dbldatagen`, `databricks-feature-engineering`, and `mlflow-skinny[databricks]>=2.5.0`. These libraries are used across multiple notebooks throughout the book.

- **Photon acceleration**: Photon is an acceleration engine that speeds up ETL and SQL workloads. Photon currently is not advantageous for standard ML modeling.
- **VM types**: There are many VMs to pick from. You can choose VMs by family from the drop-down list. You can start with a VM in the *General Purpose* group if you need more clarification or are just starting.
- **Min and max workers**: The rule of thumb is to start with a few max workers and increase the number of workers as your workload increases. Keep in mind that your cluster will autoscale for you. However, we still recommend starting smaller and growing out for only the more compute-heavy examples, such as certain notebooks in the Multilabel Image Classification deep learning project.

You now have your development environment set up. You are ready to lock down your credentials for safe use in your code.

Setting up the Databricks CLI and secrets

The Databricks CLI is the command line interface for Databricks. We recommend using the web terminal from your cluster to install Databricks CLI as well as create Databricks secrets. As we mentioned earlier in the chapter, there are other options to get access to Kaggle datasets, but we walk you through the steps to set up secrets here. Please see the documentation in *Further reading* for more details on the CLI, installation, usage, and secrets.

1. Go to the **Apps** tab of the compute cluster you set up in the last section. You can refer to *Figure 2.13* to see the location of the **Apps** tab. Apps are only available while the cluster is, so it may be greyed out initially. You will have to start your cluster to proceed.
2. Select the web terminal.
3. Install the latest version of the CLI. `curl -fsSL https://raw.githubusercontent.com/databricks/setup-cli/main/install.sh | sudo sh`
4. Check your Databricks version to be sure it's greater than `0.2`. We had to point to the updated version in the location installed by curl. `/usr/local/bin/databricks-v`.
5. Next, you need to configure the connection. You need your PAT for this:

```
/usr/local/bin/databricks configure
```

6. Create a secret scope for storing credentials related to this book:

```
/usr/local/bin/databricks secrets create-scope "machine-
learning-in-action"
```