Out of the hundred models, only six have an R-squared value of `0.85` or higher. Databricks AutoML saved us considerable time and effort! During the experiment, MLflow tried a variety of algorithms and hyperparameters, tuning with **HyperOpt**. This experiment is also distributed with the power of Spark, meaning we have a solid model that was tuned efficiently. Our best-performing model, which will become the baseline model, performs reasonably well, and indicates the dataset has predictive power. In practice, the next step might be to improve performance by creating new features, gathering more data points, or enriching the dataset. In this book, we move on to preparing for production.

Project: Streaming Transactions

You will create a `training_set` spec similar to the Favorita Sales project for the Streaming Transactions project. Before we jump in, let's remember where we are and where we are going by reviewing the project pipeline for our streaming use case.

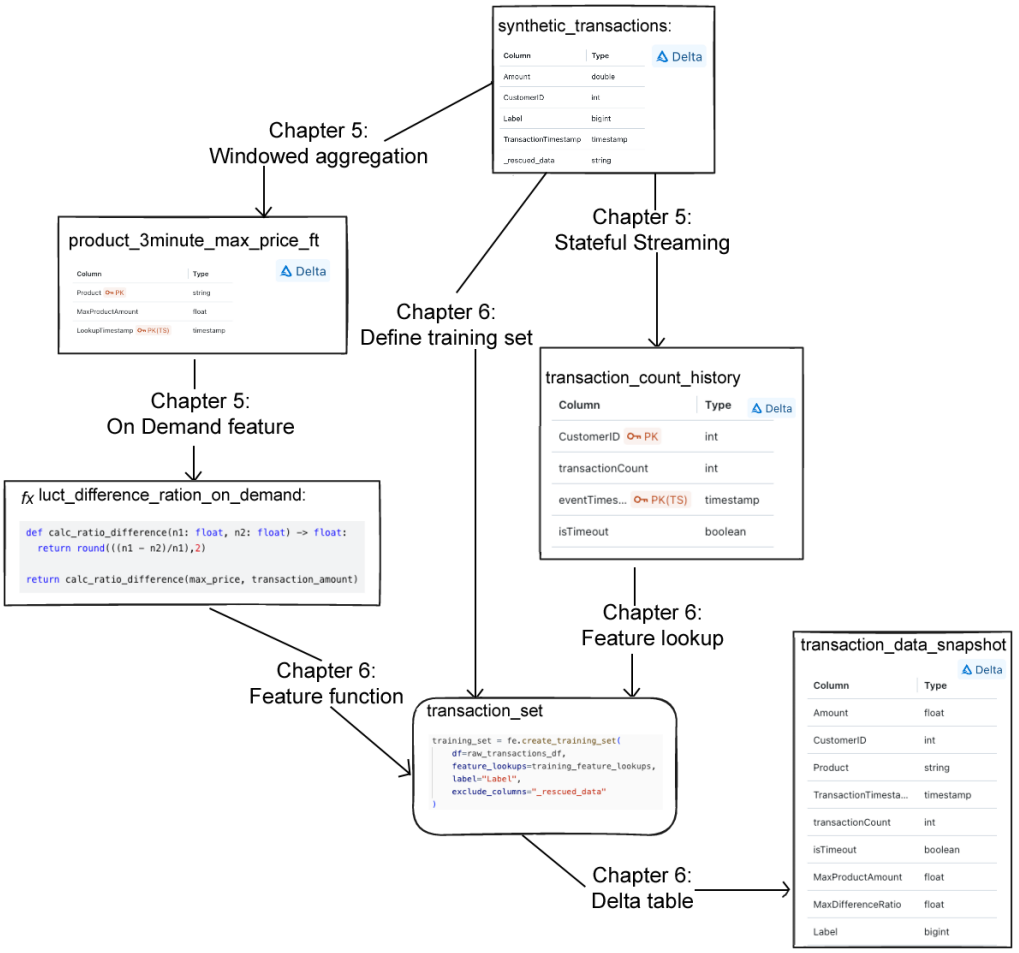


Figure 6.15 – The project pipeline for the Streaming Transactions Project