Walk through the two workflow paths, starting with the upper path. The batch features are updated, thus providing feature data for inference. The data is ready for predictions, and the inference task can begin.

The lower path has the updating of labels as the first task. It adds the latest data to the `transaction_labels` table and merges all new labels that match previous predictions into the inference table. Skip forward beyond the first iteration, the inference table contains not only previous predictions but also the labels for those predictions. The model training is performed using the updated table containing the features. The retraining of the model only occurs if there is data in the inference table, as shown in *Figure 7.27*. The retraining is, of course, followed by validation when needed. The validation notebook exits when it detects that the latest version of the model has already been tested.

```
if not spark.catalog.tableExists(table_name) or spark.table(tableName=table_name).isEmpty():
  dbutils.notebook.exit("No inference table exists yet")
```

Figure 7.27 - The retraining notebook exits if there is no data to retrain on

Create the workflow job shown in *Figure 7.29*. You can reference the configuration for the job in the `model retraining n inference workflow.yaml` file. The workflow automatically provides the lineage of all upstream and downstream tables. You can see these in *Figure 7.26*. This saves us time on documentation.

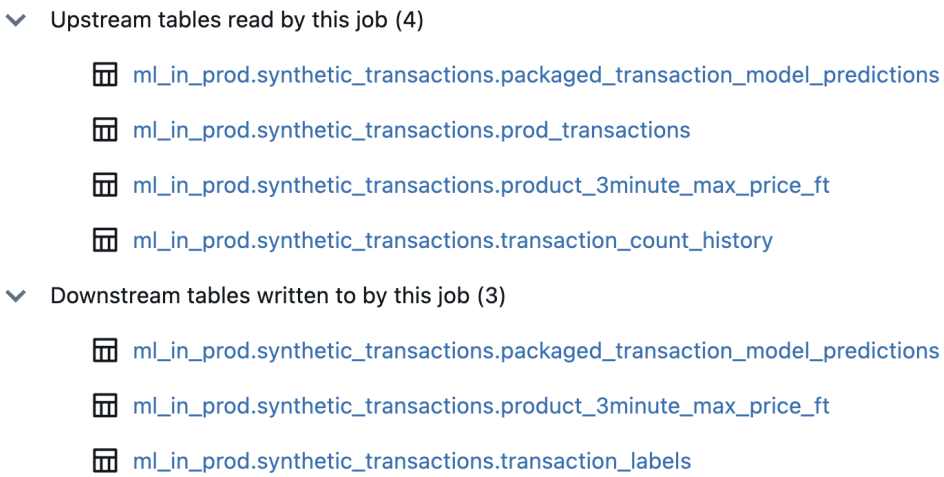


Figure 7.28 – The workflow table lineage